

In [73]:

```
import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, roc_curve, auc
from sklearn.model_selection import train_test_split
from math import exp
```

1 The Bayes Classifier

In [20]:

```
#a. Set random number generator seed
np.random.seed(2019)
```

In [21]:

```
#b. Simulate a dataset of N = 200 with X1, X2
x1 = np.random.uniform(-1, 1, 200)
x2 = np.random.uniform(-1, 1, 200)
```

In [22]:

```
#c. Calculate  $Y = X1 + X1^2 + X2 + X2^2 + \epsilon$ , where  $\epsilon \sim N(\mu=0, \sigma^2=0.25)$ .
y = x1 + x1**2 + x2 + x2**2 + np.random.normal(0, 0.5, 200)
```

In [23]:

```
#d. Y is defined in terms of the log-odds of success on the domain [-∞, +∞]. Calculate the probability of success bounded between [0, 1].
```

```
def cal_prob(x):  
    p = np.exp(np.array(x))/(1 + np.exp(np.array(x)))  
    return p
```

```
p = cal_prob(Y)  
p_tf=p>0.5  
p
```

Out[23]:

```
array([0.82849187, 0.4608178 , 0.51176518, 0.6051898  
2, 0.90938358,  
      0.59006533, 0.96321944, 0.75761982, 0.9496808  
8, 0.48865303,  
      0.52303872, 0.90609121, 0.74523342, 0.6402173  
3, 0.79025966,  
      0.96359111, 0.53661513, 0.47997726, 0.3211723  
3, 0.56016197,  
      0.5738233 , 0.56393526, 0.47938579, 0.5309953  
, 0.59683373,  
      0.54477001, 0.57012001, 0.44425388, 0.8319347  
4, 0.37961348,  
      0.69761183, 0.85921594, 0.6296643 , 0.9486115  
1, 0.84212708,  
      0.71464094, 0.6528656 , 0.63641372, 0.4685504  
2, 0.72024099,  
      0.83368541, 0.2684377 , 0.50721412, 0.8443489  
3, 0.60625281,  
      0.91224952, 0.75334057, 0.75879306, 0.5710054  
8, 0.98242394,  
      0.4034332 , 0.46750637, 0.47205267, 0.8283087  
4, 0.67958694,  
      0.84624222, 0.38998106, 0.60138008, 0.4303328  
5, 0.58575614,  
      0.72992691, 0.44086161, 0.21546613, 0.8366079  
1, 0.70091667,  
      0.65390153, 0.8829348 , 0.72034761, 0.4635184  
2, 0.85942801,  
      0.29297696, 0.40447946, 0.65573051, 0.8379782
```

3, 0.63467034,
0.7319606 , 0.97979748, 0.22591391, 0.7814395
4, 0.64677296,
0.27131764, 0.66862327, 0.60950501, 0.4806255
1, 0.93391771,
0.56902771, 0.38568622, 0.348932 , 0.6249423
8, 0.34982382,
0.35166567, 0.81298598, 0.59648047, 0.8695733
3, 0.71943754,
0.45679997, 0.48496706, 0.72286644, 0.2953835
9, 0.79582533,
0.46760114, 0.67828772, 0.67883433, 0.6516894
1, 0.83776143,
0.61611539, 0.98303505, 0.59678071, 0.6008586
3, 0.52494716,
0.53873482, 0.63684489, 0.50774389, 0.6236085
2, 0.63971866,
0.95834086, 0.70657401, 0.37470773, 0.5739381
, 0.34114826,
0.43596874, 0.80545546, 0.16948187, 0.5752372
5, 0.47504385,
0.7978173 , 0.91738198, 0.76164394, 0.8629260
7, 0.5341799 ,
0.43914779, 0.44464322, 0.3554489 , 0.2732362
3, 0.89913037,
0.81451059, 0.64782501, 0.92054441, 0.8921602
1, 0.31424726,
0.9623672 , 0.92904984, 0.9808496 , 0.6944294
7, 0.50784292,
0.48764802, 0.41643588, 0.69705644, 0.2148158
, 0.50237935,
0.69171392, 0.40122293, 0.75039878, 0.8135467
5, 0.9422711 ,
0.84681109, 0.28927443, 0.67763596, 0.3209347
2, 0.57147493,
0.39206698, 0.70836429, 0.3435361 , 0.5415438
4, 0.69181114,
0.27088757, 0.61302608, 0.88045321, 0.5409831
3, 0.85206581,
0.34219011, 0.35369037, 0.5021619 , 0.6913736
8, 0.77531435,
0.53858279, 0.35298938, 0.6533829 , 0.6715022
2, 0.81353448,
0.8856535 , 0.84202067, 0.24510023, 0.7047935
8, 0.957143 ,

```
0.44740698, 0.36650411, 0.82491493, 0.9192492
1, 0.70098476,
0.52979499, 0.24490036, 0.86720487, 0.4514611
7, 0.59474726,
0.62640475, 0.70520926, 0.23437945, 0.6720264
3, 0.90593173])
```

In [27]:


```
#e-h.
```

```
nb = GaussianNB().fit(X.transpose(), p_tf)
xx, yy = np.meshgrid(np.linspace(-1, 1, 100), np.linspace(-1, 1,
100))
Z = nb.predict_proba(np.c_[xx.ravel(), yy.ravel()])
Z = Z[:,1].reshape(xx.shape)

data = np.vstack((X, p_tf))
df = pd.DataFrame(dict(x=data[0,:], y=data[1,:], label=data[2,:])
))

groups = df.groupby('label')
f, x = plt.subplots()
for name, group in groups:
    x.plot(group.x, group.y, marker='o', linestyle='', ms=8, lab
el=name)
x.contour(xx, yy, Z, [0.5], colors='blue')

plt.xlabel('X1')
plt.ylabel('X2')
plt.title('The Bayes Classifier')
plt.legend(['success', 'failure'], loc=1)
plt.show()
```



2 Exploring Simulated Differences between LDA and QDA

2

If the Bayes boundary is linear, QDA will perform better on the training set. Because the higher the flexibility, the closer fit will QDA get. But LDA will perform better on the test set, because QDA could overfit the training set, which leads to a bad performance on the Bayes decision boundary.

In [39]:

```
#a. Repeat the following process 1000 times.
```

```
def linear(n):
```

```
    #Simulate a dataset of 1000 observation
```

```
    x1_2 = np.random.uniform(-1,1,1000)
```

```
    x2_2 = np.random.uniform(-1,1,1000)
```

```
    Y = (x1_2 + x2_2 + np.random.normal(size = 1000))>=0
```

```
    data=np.vstack((x1_2, x2_2, Y))
```

```
    s = np.arange(1000)
```

```
    np.random.shuffle(s)
```

```
    #Randomly split the dataset into 70/30% training/test se
```

ts

```
    train_idx = seq[:700]
```

```
    test_idx = seq[700:]
```

```
    train, test = data[:,train_idx], data[:,test_idx]
```

```
    x_train = train[:,2,:].T
```

```
    y_train = train[2,:].T
```

```
    x_test=test[:,2,:].T
```

```
    y_test=test[2,:].T
```

```
    #Use the training dataset to estimate LDA and QDA models
```

.

```
    # LDA model
```

```
    lda = LinearDiscriminantAnalysis()
```

```
    lda_model = lda.fit(x_train,y_train)
```

```
    # QDA model
```

```
    qda = QuadraticDiscriminantAnalysis()
```

```
    qda_model= qda.fit(x_train, y_train)
```

```
    # Calculate each model's training and test error rate.
```

```
    lda_train = 1 - lda_model.score(x_train, y_train)
```

```
    lda_test = 1 - lda_model.score(x_test, y_test)
```

```
    qda_train = 1 - qda_model.score(x_train, y_train)
```

```
    qda_test  = 1 - qda_model.score(x_test, y_test)
```

```
    return lda_train, lda_test, qda_train, qda_test
```

In [41]:

```
# Repeat the following process 1000 times.
lda_train_err = np.zeros(1000)
lda_test_err = np.zeros(1000)
qda_train_err = np.zeros(1000)
qda_test_err = np.zeros(1000)

for i in range(1000):
    lda_train, lda_test, qda_train, qda_test=linear(i)
    lda_train_err[i] = lda_train
    lda_test_err[i] = lda_test
    qda_train_err[i] = qda_train
    qda_test_err[i] = qda_test
```

In [42]:

```
#b. Summarize all the simulations' error rates and report the re
sults in tabular and graphical form.

df = pd.DataFrame({'LDA_train':lda_train_err,
                   'LDA_test':lda_test_err,
                   'QDA_train':qda_train_err,
                   'QDA_test':qda_test_err
                   })


df.describe()
```

Out[42]:

	LDA_train	LDA_test	QDA_train	QDA_test
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.273091	0.276090	0.272161	0.276303
std	0.016366	0.025906	0.016601	0.026090
min	0.222857	0.200000	0.220000	0.203333
25%	0.261429	0.260000	0.261429	0.260000
50%	0.272857	0.276667	0.272857	0.276667
75%	0.284286	0.293333	0.282857	0.293333
max	0.327143	0.360000	0.328571	0.356667

In [45]:

```
# Draw the boxplot
sns.boxplot(data=df)
plt.show()
```



According to the result, QDA performs better on the training set LDA performs better on test set. This supports my answer above.

3

If the Bayes boundary is nonlinear, QDA performs better on both the training set and test set. Because the high flexibility allows QDA to perform better on describing the nonlinear relationship.

In [53]:

```
#a. Repeat the following process 1000 times.

def nonlinear(n):
    np.random.seed(n)
    #Simulate a dataset of 1000 observation
    x1_2 = np.random.uniform(-1,1,1000)
    x2_2 = np.random.uniform(-1,1,1000)
    Y = (x1_2 + x1_2**2 + x2_2 + x2_2**2 + np.random.normal(
size = 1000))>=0
    data=np.vstack((x1_2, x2_2, Y))
    s = np.arange(1000)
    np.random.shuffle(s)

    #Randomly split your dataset into 70/30% training/test sets

    train_idx = s[:700]
    test_idx = s[700:]
    train, test = data[:,train_idx], data[:,test_idx]
    x_train = train[:2,:].T
    y_train = train[2,:].T
    x_test=test[:2,:].T
    y_test=test[2,:].T

    #Use the training dataset to estimate LDA and QDA models
    .

    # LDA model
    lda = LinearDiscriminantAnalysis()
    lda_model = lda.fit(x_train,y_train)
    # QDA model
    qda = QuadraticDiscriminantAnalysis()
    qda_model= qda.fit(x_train, y_train)

    # Calculate each model's training and test error rate.
    lda_train = 1 - lda_model.score(x_train, y_train)
```

```
lda_test = 1 - lda_model.score(x_test, y_test)
qda_train = 1 - qda_model.score(x_train, y_train)
qda_test = 1 - qda_model.score(x_test, y_test)

return lda_train, lda_test, qda_train, qda_test
```

```
lda_train_err = np.zeros(1000)
lda_test_err = np.zeros(1000)
qda_train_err = np.zeros(1000)
qda_test_err = np.zeros(1000)
```

```
for i in range(1000):
    lda_train, lda_test, qda_train, qda_test=nonlinear(i)
    lda_train_err[i] = lda_train
    lda_test_err[i] = lda_test
    qda_train_err[i] = qda_train
    qda_test_err[i] = qda_test
```

In [50]:

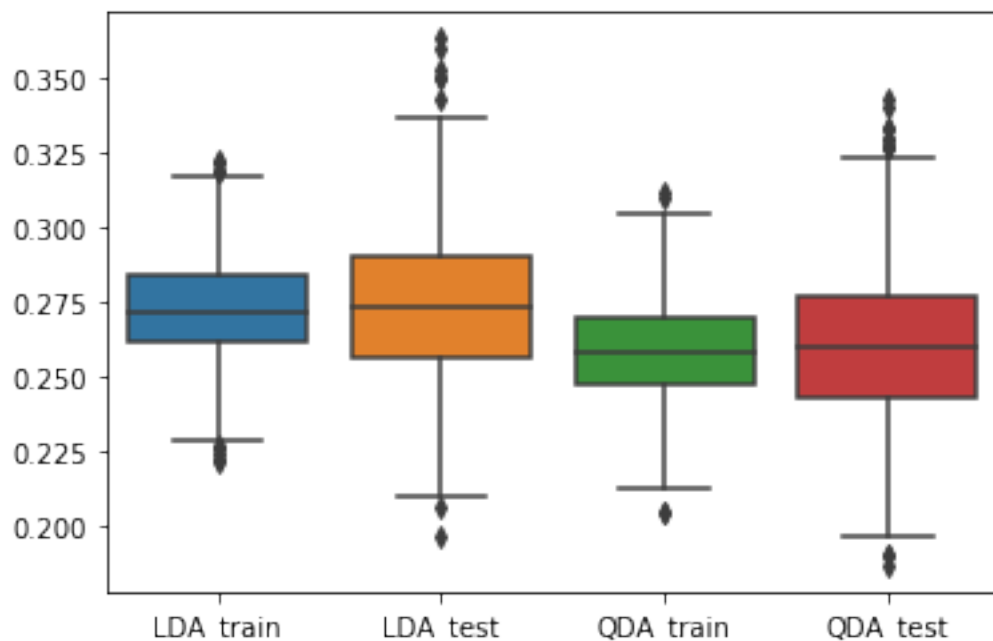
```
#b. Summarize all the simulations' error rates and report the re  
sults in tabular and graphical form.  
df = pd.DataFrame({'LDA_train':lda_train_err,  
                   'LDA_test':lda_test_err,  
                   'QDA_train':qda_train_err,  
                   'QDA_test':qda_test_err  
                   })  
#Summarize all the simulations' error rates  
df.describe()
```

Out[50]:

	LDA_train	LDA_test	QDA_train	QDA_test
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.272387	0.274833	0.259020	0.261623
std	0.016924	0.025222	0.016340	0.024093
min	0.221429	0.196667	0.204286	0.186667
25%	0.261429	0.256667	0.247143	0.243333
50%	0.271429	0.273333	0.258571	0.260000
75%	0.284286	0.290000	0.270000	0.276667
max	0.322857	0.363333	0.311429	0.343333

In [52]:

```
# Draw the boxplot
sns.boxplot(data=df)
plt.show()
```



According to the result, QDA performs better on both the training set and test set. This supports my answer above.

4

When the sample size is relatively small, the test prediction accuracy of LDA is better than QDA. Because LDA is less likely to be overfitting than QDA. When sample size is large enough, the test prediction accuracy of QDA is better than LDA. Because the higher the flexibility, the closer fit will QDA gets and variance won't be a problem in terms of the larger sample sizes.

In [93]:

```
np.random.seed(10)
def nonlinear_new(m):
    #Simulate a dataset of 1000 observation
    x1_3 = np.random.uniform(-1,1,m)
    x2_3 = np.random.uniform(-1,1,m)
    Y = (x1_3 + x1_3**2 + x2_3 + x2_3**2 + np.random.normal(
size = m))>=0
```

```

size = m) >= 0
    data=np.vstack((x1_3, x2_3, Y))

    s = np.arange(m)
    np.random.shuffle(s)
    #Randomly split your dataset into 70/30% training/test sets

    train_idx = s[:int(0.7*m)]
    test_idx = s[int(0.7*m):]
    train, test = data[:,train_idx], data[:,test_idx]
    x_train = train[:2,:].T
    y_train = train[2,:].T
    x_test=test[:2,:].T
    y_test=test[2,:].T

    #Use the training dataset to estimate LDA and QDA models
    .

    # LDA model
    lda = LinearDiscriminantAnalysis()
    lda_model = lda.fit(x_train,y_train)
    # QDA model
    qda = QuadraticDiscriminantAnalysis()
    qda_model= qda.fit(x_train, y_train)

    # Calculate each model's training and test error rate.
    lda_train = 1 - lda_model.score(x_train, y_train)
    lda_test = 1 - lda_model.score(x_test, y_test)
    qda_train = 1 - qda_model.score(x_train, y_train)
    qda_test  = 1 - qda_model.score(x_test, y_test)

    return lda_train, lda_test, qda_train, qda_test

# simulating 1000 times for N = [100, 1000, 10000, 100000]
N = [10**2, 10**3, 10**4, 10**5]
lda_err = np.zeros(len(N))
qda_err = np.zeros(len(N))
for i in range(len(N)):
    lda_test_err = np.zeros(1000)
    qda_test_err = np.zeros(1000)

    for j in range(1000):
        lda_train, lda_test, qda_train, qda_test= nonlinear_new(
N[i])
        lda_test_err[j] = lda_test
        qda_test_err[j] = qda_test

```

```
lda_err[i] = lda_test_err.mean()  
qda_err[i] =qda_test_err.mean()
```

In [94]:


```
d = {'Sample Size':['100', '1000', '10000', '100000'],  
     'LDA_error':lda_err,  
     'QDA_error':qda_err}  
  
df = pd.DataFrame(d)  
df
```

Out[94]:

	Sample Size	LDA_error	QDA_error
0	100	0.286333	0.274367
1	1000	0.273327	0.260773
2	10000	0.273855	0.261025
3	100000	0.273349	0.260569

In [97]:

```
ratio=np.array(qda_err/lda_err)
plt.plot([1,2,3,4], ratio)
plt.xticks([1,2,3,4], ['100', '1000', '10000', '100000'])
plt.xlabel('Number of observations')
plt.ylabel('QDA/LDA test error ratio')
plt.title('QDA/LDA Test error ratio vs Number of observations')
plt.show()
```



According to the above graph, we can expect the test error rate of QDA relative to LDA to decline, which proves my answer above.

5 Modeling voter turnout

In [64]:

```
mh = pd.read_csv('mental_health.csv')
mh.dropna(inplace=True)
mh.head()
```

Out[64]:

	vote96	mhealth_sum	age	educ	black	female	married	inc10
0	1.0	0.0	60.0	12.0	0	0	0.0	4.8149
2	1.0	1.0	36.0	12.0	0	0	1.0	8.8273
3	0.0	7.0	21.0	13.0	0	0	0.0	1.7387
7	0.0	6.0	29.0	13.0	0	0	0.0	10.6998
11	1.0	1.0	41.0	15.0	1	1	1.0	8.8273

In [75]:

```
#a. Split the data into a training and test set (70/30)
np.random.seed(124)
vote96 = mh['vote96']
x = mh.drop(columns=['vote96'])
x_train, x_test, y_train, y_test = train_test_split(x, vote96, t
est_size=0.3)
```


In [76]:

```
#b. Using the training set and all important predictors, estimate the following models with vote96 as the response variable:
models=[]
#i. Logistic regression model
logit = LogisticRegression().fit(x_train, y_train)
models.append(('Logistic', logit))

#ii. Linear discriminant model
lda = LinearDiscriminantAnalysis().fit(x_train, y_train)
models.append(('LDA', lda))

#iii. Quadratic discriminant model
qda = QuadraticDiscriminantAnalysis().fit(x_train, y_train)
models.append(('QDA', qda))

#iv. Naive Bayes (you can use the default hyperparameter settings)
nb = GaussianNB().fit(x_train, y_train)
models.append(('Naive Bayes', nb))

#v. K-nearest neighbors with K = 1,2,...,10 (that is, 10 separate models varying K) and Euclidean distance metrics
m_knn = []
for k in range(1,11):
    knn= KNeighborsClassifier(n_neighbors=k).fit(x_train, y_train)
    models.append(('KNN_{}'.format(k), knn))
```

```
//anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
```

```
FutureWarning)
```

In [92]:

```
#c. Using the test set, create the list for prediction, names of  
the models, prediction_probability,  
preds = []  
names=[]  
probs = []  
  
for name, model in models:  
    preds.append(model.predict(x_test))  
    names.append(name)  
    probs.append(model.predict_proba(x_test)[:, 1])  
  
error_rates = [ (1 - accuracy_score(y_test, pred)) for pred in p  
reds]  
  
dict = {'Model Name':names,  
        'Error rate': error_rates}  
  
df = pd.DataFrame(dict)  
df
```

Out[92] :

	Model Name	Error rate
0	Logistic	0.271429
1	LDA	0.277143
2	QDA	0.300000
3	Naive Bayes	0.285714
4	KNN_1	0.345714
5	KNN_2	0.371429
6	KNN_3	0.322857
7	KNN_4	0.334286
8	KNN_5	0.300000
9	KNN_6	0.308571
10	KNN_7	0.291429
11	KNN_8	0.288571
12	KNN_9	0.274286
13	KNN_10	0.274286

In [113]:

```
def auc_roc(prob, name):
    fpr, tpr, _ = roc_curve(y_test, prob)
    aucs=auc(fpr, tpr)
    print(name+ ': AUC = %f' % (aucs))
# plot the roc curve for the model
    plt.plot(fpr, tpr, label=name)
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.legend()

plt.figure(figsize=(12,10))
rand_probs = [0] * len(y_test)
rand_fpr, rand_tpr, _ = roc_curve(y_test, rand_probs)
plt.plot(rand_fpr, rand_tpr, linestyle='--', label='Random Class
ifier')
for i, model in enumerate(models):
    auc_roc(probs[i], names[i])
```

Logistic: AUC = 0.757322

LDA: AUC = 0.756430

QDA: AUC = 0.736247

Naive Bayes: AUC = 0.745503

KNN_1: AUC = 0.614407

KNN_2: AUC = 0.660608

KNN_3: AUC = 0.685753

KNN_4: AUC = 0.695863

KNN_5: AUC = 0.715265


KNN_6: AUC = 0.726862

KNN_7: AUC = 0.744555

KNN_8: AUC = 0.749350

KNN_9: AUC = 0.747863

KNN_10: AUC = 0.761095



d. According to the above analyses, the best models for this dataset are Logistic and LDA, which produces the highest AUC and lowest error rate.