

# HW02 - Classification [MACS 30100]

Adarsh Mathew

02/01/2020

## Q1 - Naive Bayes Decision Boundary

```
# Generating the table

set.seed(01302020)
n_obs <- 200

x_1 <- runif(n_obs, -1, 1)
x_2 <- runif(n_obs, -1, 1)
err <- rnorm(n_obs, mean = 0, sd = 0.5)

nb_df <- cbind(x_1, x_2, err) %>%
  as_tibble() %>%
  mutate(f_x = x_1 + x_1^2 + x_2 + x_2^2 + err,
         y = exp(f_x)/(1+exp(f_x)),
         y_class = as.factor(y > 0.5))

head(nb_df)

## # A tibble: 6 x 6
##       x_1      x_2      err    f_x      y y_class
##   <dbl>   <dbl>   <dbl>  <dbl> <dbl> <fct>
## 1 -0.985 -0.0760  0.221   0.137 0.534 TRUE
## 2 -0.277  0.755   0.0937  1.22  0.772 TRUE
## 3  0.730  0.464  -0.0844  1.86  0.865 TRUE
## 4 -0.477  0.0954 -0.397  -0.542 0.368 FALSE
## 5  0.0770  0.466  -0.871  -0.106 0.474 FALSE
## 6 -0.0893 -0.447   0.00252 -0.326 0.419 FALSE

# Training the Naive Bayes model

nb_mod <- train(x = nb_df %>% select(x_1, x_2), y = nb_df$y_class, method = "nb")

confusionMatrix(nb_mod)

## Bootstrapped (25 reps) Confusion Matrix
##
## (entries are percentual average cell counts across resamples)
##
##           Reference
## Prediction FALSE TRUE
##      FALSE 19.6  7.8
```

```
##      TRUE      8.6 64.0
##
## Accuracy (average) : 0.836
# Setting up the grid-search to plot the decision boundary

nb_grid <- expand.grid(
  x_1 = seq(min(nb_df$x_1), max(nb_df$x_1), length = n_obs),
  x_2 = seq(min(nb_df$x_2), max(nb_df$x_2), length = n_obs)
)

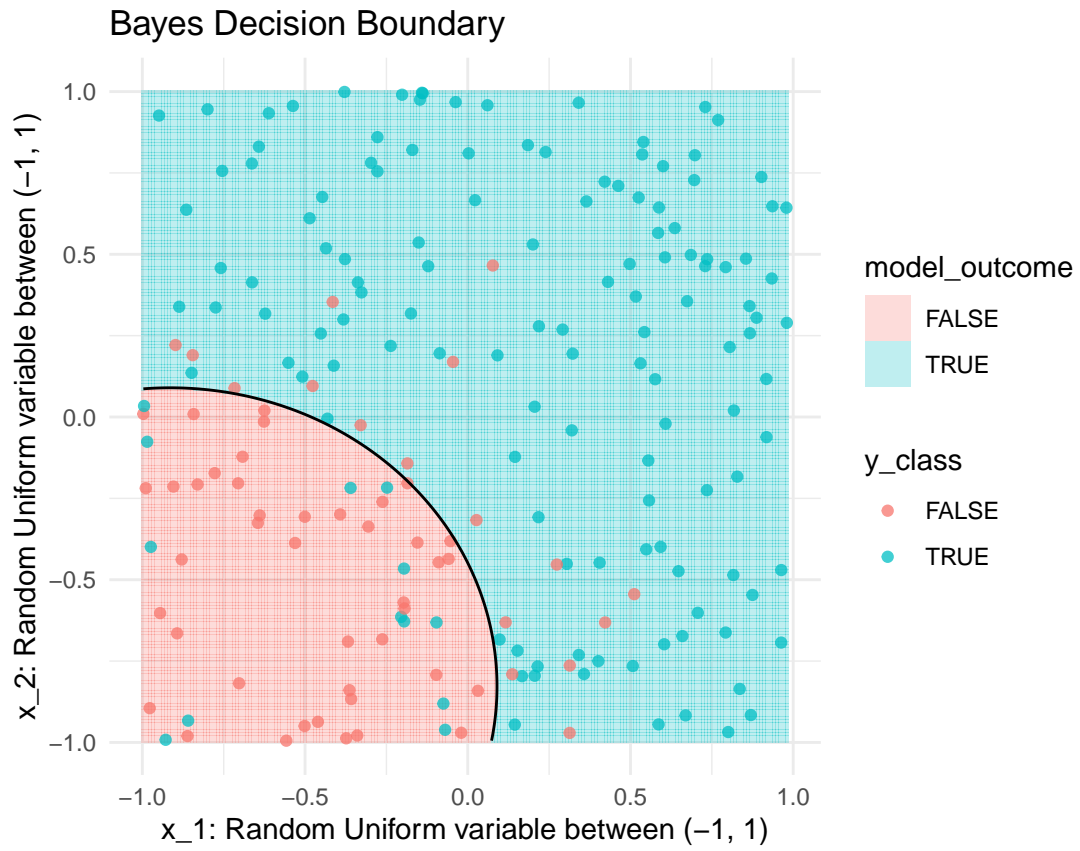
nb_grid <- nb_grid %>%
  bind_cols('model_outcome' = predict(nb_mod, newdata = nb_grid),
            'grid_prob' = predict(nb_mod, newdata = nb_grid, type = 'prob')) %>%
  rename(prob_true = 'TRUE', prob_false = 'FALSE')

head(nb_grid)

##      x_1      x_2 model_outcome prob_false prob_true
## 1 -0.9968940 -0.9945238      FALSE  0.8320680 0.1679320
## 2 -0.9869575 -0.9945238      FALSE  0.8324237 0.1675763
## 3 -0.9770210 -0.9945238      FALSE  0.8327334 0.1672666
## 4 -0.9670845 -0.9945238      FALSE  0.8329972 0.1670028
## 5 -0.9571480 -0.9945238      FALSE  0.8332153 0.1667847
## 6 -0.9472115 -0.9945238      FALSE  0.8333879 0.1666121

# Plotting the decision boundary

nb_df %>% ggplot(aes(x = x_1, y = x_2)) +
  geom_tile(data = nb_grid, aes(fill = model_outcome), alpha = 0.25) +
  geom_point(aes(color = y_class), alpha = .75) +
  geom_contour(data = nb_grid, aes(z = prob_true), colour = "black", breaks = .5) +
  theme_minimal() +
  #theme(legend.position = "top") +
  coord_equal() +
  ggtitle("Bayes Decision Boundary") +
  labs(x = "x_1: Random Uniform variable between (-1, 1)",
       y = "x_2: Random Uniform variable between (-1, 1)")
```



# LDA-vs-QDA

Adarsh Mathew

2/1/2020

## Q2: LDA vs QDA with a linear Decision Boundary

```
model_df_gen <- function(seed_val, n_obs, sec_order_terms, split_prop, return_type){  
  
  set.seed(seed_val)  
  
  x_1 <- runif(n_obs, -1, 1)  
  x_2 <- runif(n_obs, -1, 1)  
  err <- rnorm(n_obs, mean = 0, sd = 1)  
  
  df <- cbind(x_1, x_2, err) %>%  
    as_tibble(rownames = NULL) %>%  
    mutate(y = ifelse(sec_order_terms == TRUE, x_1 + x_2 + x_1^2 + x_2^2, x_1 + x_2),  
           y_class = factor(y >= 0, levels = c(TRUE, FALSE)),  
           f_x = y + err,  
           y_sim_class = factor(f_x >= 0, levels = c(TRUE, FALSE)))  
  
  df_split <- rsample::initial_split(df, prop = split_prop)  
  df_train <- rsample::training(df_split) %>% select(x_1, x_2, y_sim_class)  
  df_test <- rsample::testing(df_split) %>% select(x_1, x_2, y_sim_class)  
  
  if(return_type == "train") return(df_train) else return(df_test)  
}  
  
model_runner <- function(df_train, model_type){  
  df_train_ed <- droplevels(df_train)  
  
  mod <- train(  
    df_train_ed %>% select("x_1", "x_2") %>% na.omit(),  
    (df_train_ed %>% na.omit())$y_sim_class,  
    metric = 'Accuracy',  
    method = model_type,  
    allowParallel = FALSE  
  )  
  
  return(mod)  
}  
  
model_acc_gen <- function(model_obj, pred_df){  
  
  pred_df_aug <- pred_df %>%
```

```

    bind_cols('model_outcome' = predict(model_obj, newdata = pred_df),
              'model_prob' = predict(model_obj, newdata = pred_df, type = 'prob')['TRUE']) %>%
    rename(prob_true = 'TRUE')

    #print("Training Accuracy:")
    pred_acc <- postResample(pred_df_aug$model_outcome, pred_df$y_sim_class)[1]

    return(as.numeric(pred_acc))
}

```

```

n_iter <- 1000
#n_cores <- availableCores() - 2
#plan(multicore, workers = n_cores)

discrim_df <- sample(1:1000, n_iter, replace=T) %>%
  as_tibble(rownames = NULL) %>%
  rename(seed_value = value) %>%
  mutate(train_df = map(seed_value, ~model_df_gen(.x, 1000, FALSE, 0.7, "train")),
         test_df = map(seed_value, ~model_df_gen(.x, 1000, FALSE, 0.7, "test")),
         lda_mod = map(train_df, ~model_runner(.x, "lda")),
         lda_err_train = 1 - map2_dbl(lda_mod, train_df, model_acc_gen),
         lda_err_test = 1 - map2_dbl(lda_mod, test_df, model_acc_gen),
         qda_mod = map(train_df, ~model_runner(.x, "qda")),
         qda_err_train = 1 - map2_dbl(qda_mod, train_df, model_acc_gen),
         qda_err_test = 1 - map2_dbl(qda_mod, test_df, model_acc_gen))

#future::plan(future::sequential)

head(discrim_df)

```

```

## # A tibble: 6 x 9
##   seed_value train_df test_df lda_mod lda_err_train lda_err_test qda_mod
##   <int> <list> <list> <list> <dbl> <dbl> <list>
## 1     595 <tibble~ <tibbl~ <train> 0.389 0.433 <train>
## 2     992 <tibble~ <tibbl~ <train> 0.429 0.437 <train>
## 3      39 <tibble~ <tibbl~ <train> 0.166 0.137 <train>
## 4       9 <tibble~ <tibbl~ <train> 0.109 0.103 <train>
## 5      36 <tibble~ <tibbl~ <train> 0.309 0.363 <train>
## 6     570 <tibble~ <tibbl~ <train> 0.314 0.29 <train>
## # ... with 2 more variables: qda_err_train <dbl>, qda_err_test <dbl>

```

```

discrim_df_plot <- discrim_df %>%
  pivot_longer(c('lda_err_train', 'lda_err_test', 'qda_err_train', 'qda_err_test'), names_to = "mod_dfname",
               separate(`mod_dfname`, sep = "_", into = c("model_type", "metric", "df_type")))

head(discrim_df_plot)

```

```

## # A tibble: 6 x 9
##   seed_value train_df test_df lda_mod qda_mod model_type metric df_type
##   <int> <list> <list> <list> <list> <chr> <chr> <chr>
## 1     595 <tibble~ <tibbl~ <train> <train> lda     err   train
## 2     595 <tibble~ <tibbl~ <train> <train> lda     err   test
## 3     595 <tibble~ <tibbl~ <train> <train> qda     err   train
## 4     595 <tibble~ <tibbl~ <train> <train> qda     err   test
## 5     992 <tibble~ <tibbl~ <train> <train> lda     err   train

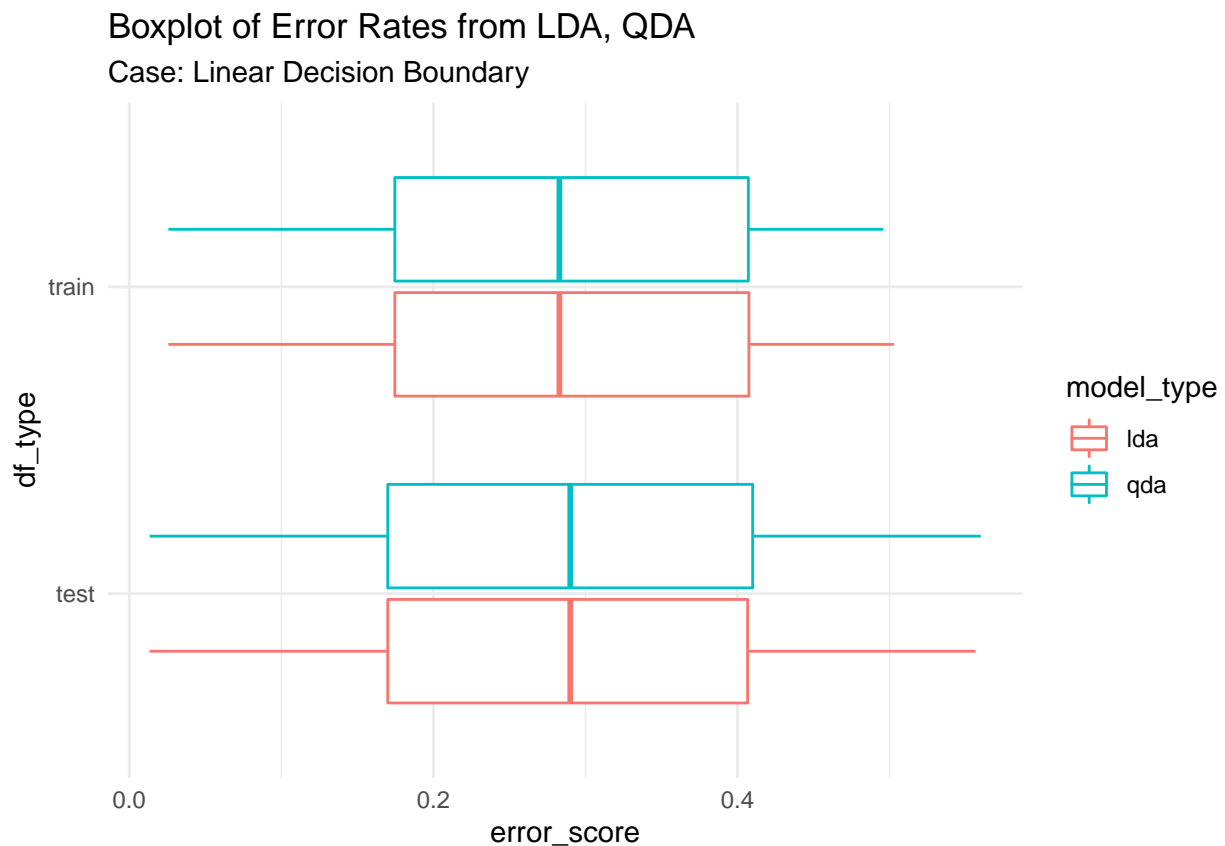
```

```
## 6          992 <tibble~ <tibble~ <train> <train> lda          err      test
## # ... with 1 more variable: error_score <dbl>

discrim_df_plot %>% group_by(model_type, df_type) %>%
  summarise(mean_error_rate = mean(error_score)) %>%
  pivot_wider(names_from = model_type, values_from = mean_error_rate)

## # A tibble: 2 x 3
##   df_type  lda    qda
##   <chr>    <dbl> <dbl>
## 1 test    0.284 0.285
## 2 train    0.279 0.279

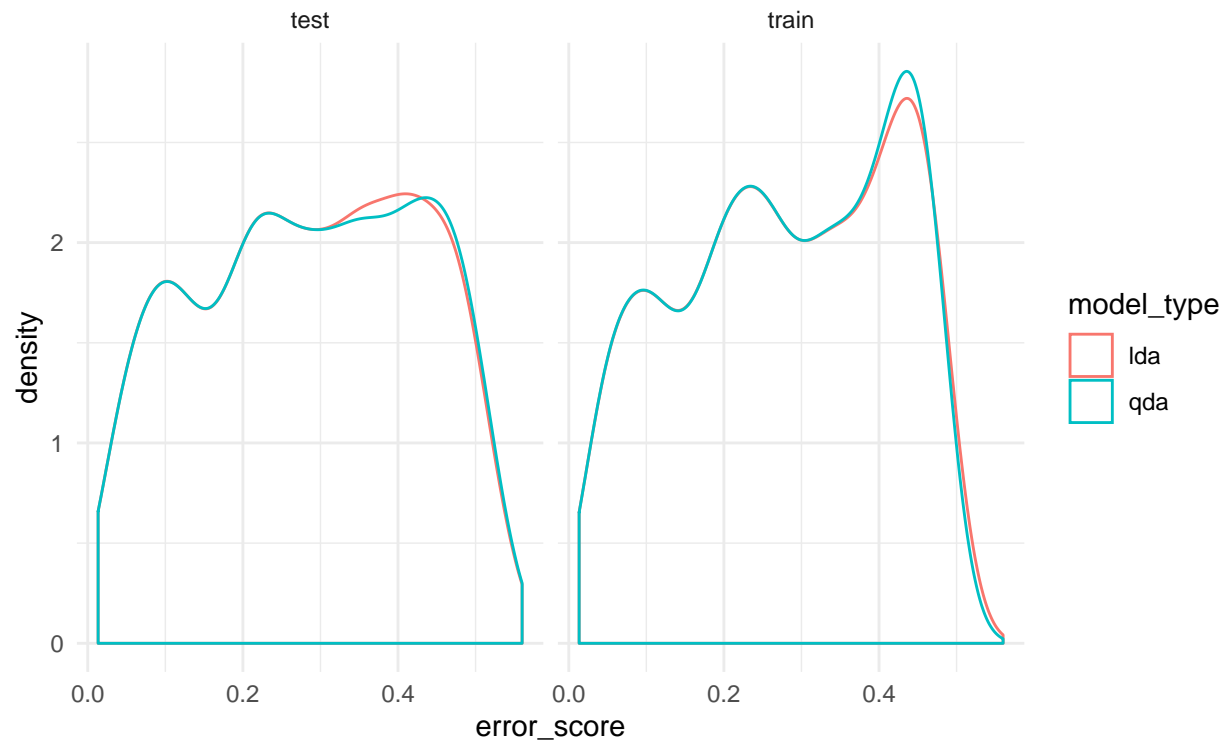
discrim_df_plot %>%
  ggplot() +
  geom_boxplot(aes(x = df_type, y = error_score, colour = model_type)) +
  theme_minimal() + coord_flip() +
  ggtitle("Boxplot of Error Rates from LDA, QDA", subtitle = "Case: Linear Decision Boundary")
```



```
discrim_df_plot %>%
  ggplot() +
  geom_density(aes(error_score, colour = model_type)) + facet_wrap(~df_type, ncol = 2) +
  theme_minimal() +
  ggtitle("Density Distribution of Error Rates from LDA, QDA", subtitle = "Case: Linear Decision Boundary")
```

## Density Distribution of Error Rates from LDA, QDA

### Case: Linear Decision Boundary



**Conclusion:** Both LDA and QDA perform about the same.

### Q3: LDA vs QDA with a non-linear Decision Boundary

```
#n_cores <- availableCores() - 2
#plan(multicore, workers = n_cores)

# discrim_df_sec <- seq(1, 1000, by = 1) %>%
#   as_tibble(rownames = NULL) %>%
#   rename(seed_value = value) %>%
#   mutate(train_df = map(seed_value, ~model_df_gen(.x, 1000, TRUE, 0.7, "train")),
#           test_df = map(seed_value, ~model_df_gen(.x, 1000, TRUE, 0.7, "test")),
#           lda_mod = map(train_df, ~model_runner(.x, "lda")),
#           lda_err_train = 1 - map2_dbl(lda_mod, train_df, model_acc_gen),
#           lda_err_test = 1 - map2_dbl(lda_mod, test_df, model_acc_gen),
#           qda_mod = map(train_df, ~model_runner(.x, "qda")),
#           qda_err_train = 1 - map2_dbl(qda_mod, train_df, model_acc_gen),
#           qda_err_test = 1 - map2_dbl(qda_mod, test_df, model_acc_gen))

#head(discrim_df_sec)
```

Since the functional programming approach is giving me inscrutable errors, let's brute-force this with a for-loop.

```
model_err_gen2 <- function(n_obs, split_prop){
```

```

#set.seed(seed_val)

#print(iter_no)

#x_1 <- runif(n_obs, -1, 1)
#x_2 <- runif(n_obs, -1, 1)
#err <- rnorm(n_obs, mean = 0, sd = 1)

df <- tibble(x_1 = runif(n_obs, -1, 1),
             x_2 = runif(n_obs, -1, 1),
             err = rnorm(n_obs, mean = 0, sd = 1))%>%
  mutate(y = x_1 + x_2 + x_1^2 + x_2^2,
         y_class = factor(y >= 0, levels = c(TRUE, FALSE)),
         f_x = y + err,
         y_sim_class = factor(f_x >= 0, levels = c(TRUE, FALSE)))

df_split <- rsample::initial_split(df, prop = split_prop)
df_train <- rsample::training(df_split) %>% select(x_1, x_2, y_sim_class) %>% na.omit()
df_test <- rsample::testing(df_split) %>% select(x_1, x_2, y_sim_class) %>% na.omit()

#df_train_ed <- droplevels(df_train)
#print("Here")

#print(head(df_train))
#print(head(df_train$y_sim_class))

# mod <- train(
#   df_train,
#   df_train$y_sim_class,
#   metric = 'Accuracy',
#   method = model_type
# )

lda_mod <- MASS::lda(y_sim_class ~ x_1 + x_2, data = df_train)
qda_mod <- MASS::qda(y_sim_class ~ x_1 + x_2, data = df_train)

train_df_aug <- df_train %>%
  bind_cols('lda_model_outcome' = predict(lda_mod, newdata = df_train)$class,
           'qda_model_outcome' = predict(qda_mod, newdata = df_train)$class)

#print("Training Accuracy:")
lda_train_err <- 1 - postResample(train_df_aug$lda_model_outcome, df_train$y_sim_class)[1]
qda_train_err <- 1 - postResample(train_df_aug$qda_model_outcome, df_train$y_sim_class)[1]

pred_df_aug <- df_test %>%
  bind_cols('lda_model_outcome' = predict(lda_mod, newdata = df_test)$class,
           'qda_model_outcome' = predict(qda_mod, newdata = df_test)$class)

#print("Training Accuracy:")

```



```
lda_test_err <- 1 - postResample(pred_df_aug$lda_model_outcome, df_test$y_sim_class)[1]
qda_test_err <- 1 - postResample(pred_df_aug$qda_model_outcome, df_test$y_sim_class)[1]

endval <- tibble("lda_train_err" = lda_train_err,
                 "lda_test_err" = lda_test_err,
                 "qda_train_err" = qda_train_err,
                 "qda_test_err" = qda_test_err)

return(endval)
}
```

```
discrim_df_sec2 <- model_err_gen2(1000, 0.7)
for(i in 1:999){
  discrim_df_sec2 <- bind_rows(discrim_df_sec2, model_err_gen2(1000, 0.7))
}
head(discrim_df_sec2)
```

```
## # A tibble: 6 x 4
##   lda_train_err lda_test_err qda_train_err qda_test_err
##   <dbl>         <dbl>         <dbl>         <dbl>
## 1      0.263      0.297      0.249      0.283
## 2      0.263      0.28      0.253      0.26
## 3      0.269      0.263      0.259      0.237
## 4      0.254      0.323      0.234      0.293
## 5      0.254      0.253      0.244      0.223
## 6      0.234      0.283      0.220      0.267
```

```
sec_discrim_df_plot <- discrim_df_sec2 %>%
  pivot_longer(c('lda_train_err', 'lda_test_err', 'qda_train_err', 'qda_test_err'), names_to = "mod_dfname",
               separate(`mod_dfname`, sep = "_", into = c("model_type", "df_type", "metric")))
head(sec_discrim_df_plot)
```

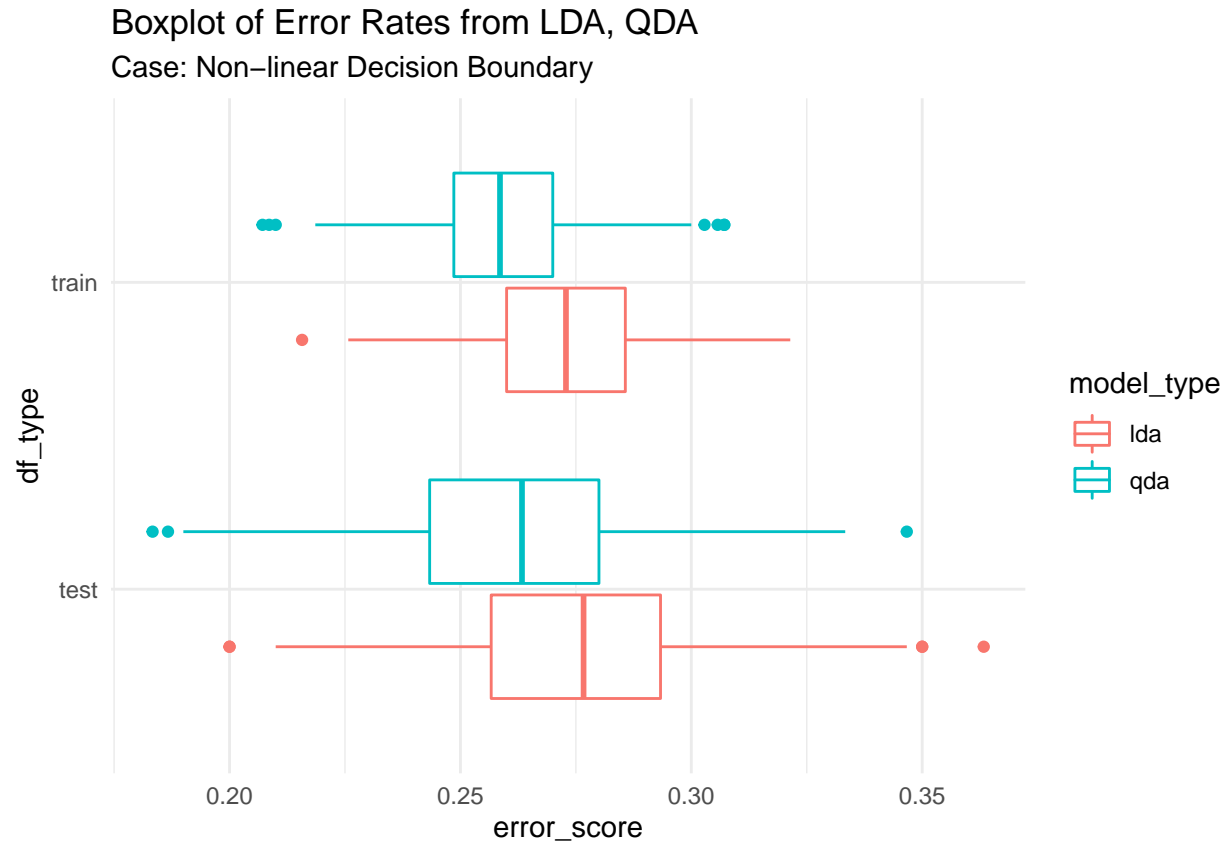
```
## # A tibble: 6 x 4
##   model_type df_type metric error_score
##   <chr>      <chr>  <chr>      <dbl>
## 1 lda      train   err        0.263
## 2 lda      test    err        0.297
## 3 qda      train   err        0.249
## 4 qda      test    err        0.283
## 5 lda      train   err        0.263
## 6 lda      test    err        0.28
```

```
sec_discrim_df_plot %>% group_by(model_type, df_type) %>%
  summarise(mean_error_rate = mean(error_score)) %>%
  pivot_wider(names_from = model_type, values_from = mean_error_rate)
```

```
## # A tibble: 2 x 3
##   df_type  lda  qda
##   <chr>   <dbl> <dbl>
## 1 test   0.276 0.262
## 2 train  0.273 0.259
```

```
sec_discrim_df_plot %>%
  ggplot() +
  geom_boxplot(aes(x = df_type, y = error_score, colour = model_type)) +
```

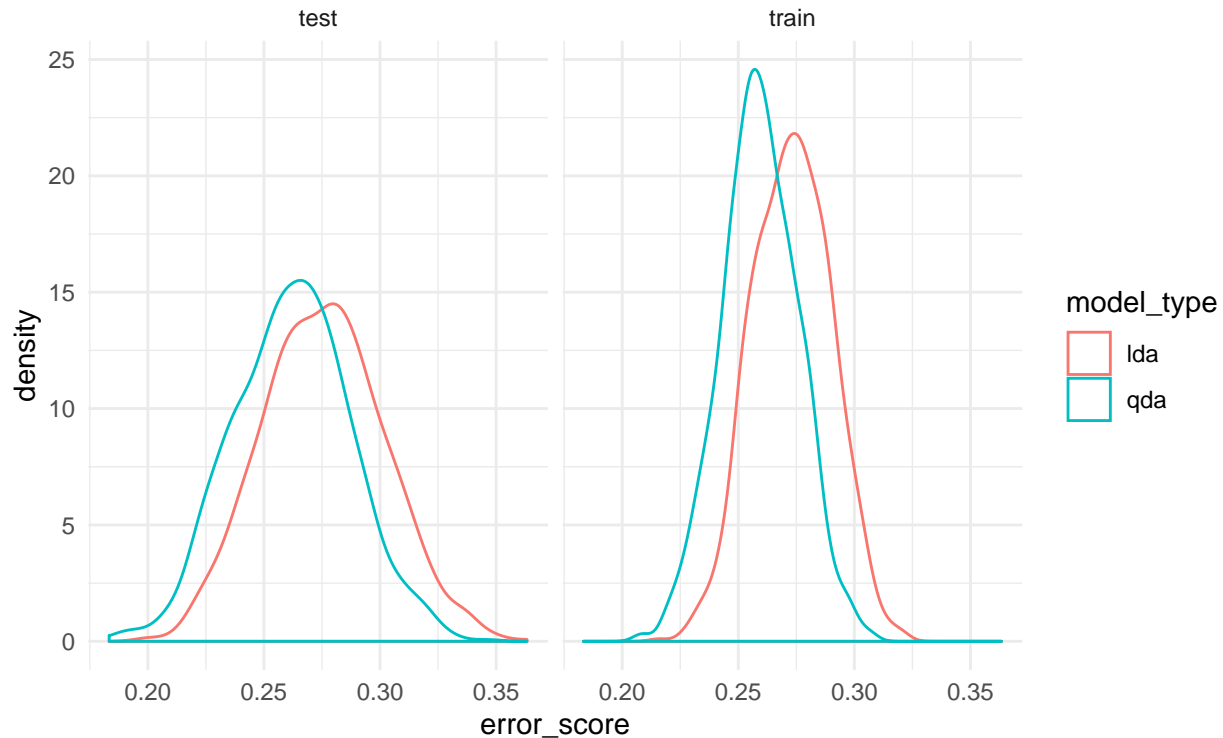
```
theme_minimal() + coord_flip() +
ggtitle("Boxplot of Error Rates from LDA, QDA", subtitle = "Case: Non-linear Decision Boundary")
```



```
sec_discrim_df_plot %>%
ggplot() +
geom_density(aes(error_score, colour = model_type)) + facet_wrap(~df_type, ncol = 2) +
theme_minimal() +
ggtitle("Density Distribution of Error Rates from LDA, QDA", subtitle = "Case: Non-linear Decision Boundary")
```

## Density Distribution of Error Rates from LDA, QDA

### Case: Non-linear Decision Boundary



**Conclusion:** We can see that QDA performs better in the case of the non-linear decision boundary.

### Q5: Effect of sample size on error rates for the non-linear decision boundary

Re-using the code from Q4:

```
n_obs_arr <- c(100,1000, 10000, 100000)

discrim_df_sec_100 <- model_err_gen2(n_obs_arr[1], 0.7)
discrim_df_sec_100 <- discrim_df_sec_100 %>% bind_cols("n_obs" = n_obs_arr[1])
for(i in 1:999){
  discrim_df_sec_100 <- bind_rows(discrim_df_sec_100, model_err_gen2(n_obs_arr[1], 0.7)%>% bind_cols("n_obs" = n_obs_arr[1]))
}

discrim_df_sec_1000 <- model_err_gen2(n_obs_arr[2], 0.7)
discrim_df_sec_1000 <- discrim_df_sec_1000 %>% bind_cols("n_obs" = n_obs_arr[2])
for(i in 1:999){
  discrim_df_sec_1000 <- bind_rows(discrim_df_sec_1000, model_err_gen2(n_obs_arr[2], 0.7)%>% bind_cols("n_obs" = n_obs_arr[2]))
}

discrim_df_sec_10000 <- model_err_gen2(n_obs_arr[3], 0.7)
discrim_df_sec_10000 <- discrim_df_sec_10000 %>% bind_cols("n_obs" = n_obs_arr[3])
for(i in 1:999){
  discrim_df_sec_10000 <- bind_rows(discrim_df_sec_10000, model_err_gen2(n_obs_arr[3], 0.7)%>% bind_cols("n_obs" = n_obs_arr[3]))
}

discrim_df_sec_100000 <- model_err_gen2(n_obs_arr[4], 0.7)
```

```

discrim_df_sec_100000 <- discrim_df_sec_100000 %>% bind_cols("n_obs" = n_obs_arr[4])
for(i in 1:999){
  discrim_df_sec_100000 <- bind_rows(discrim_df_sec_100000, model_err_gen2(n_obs_arr[4], 0.7)%>% bind_cols("n_obs" = n_obs_arr[4])
}

sim_discrim_df_sec <- bind_rows(discrim_df_sec_100, discrim_df_sec_1000, discrim_df_sec_10000, discrim_df_sec_100000,
  pivot_longer(c('lda_train_err', 'lda_test_err', 'qda_train_err', 'qda_test_err'), names_to = "mod_dfname",
    separate(`mod_dfname`, sep = "_", into = c("model_type", "df_type", "metric"))) %>%
  mutate(n_obs = as.factor(n_obs))

#sim_discrim_df_sec

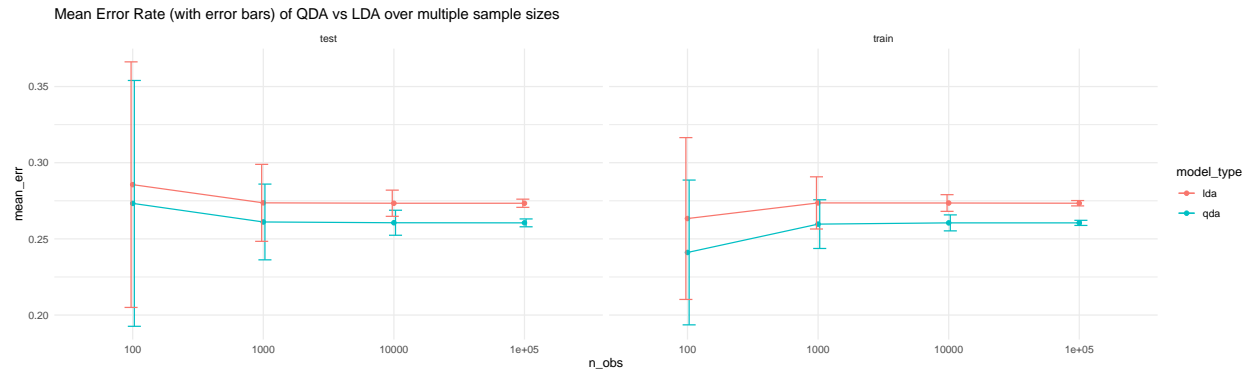
sim_discrim_df_sec_summ <- sim_discrim_df_sec %>%
  group_by(model_type, df_type, n_obs) %>%
  summarize(mean_err = mean(error_score),
    sd_err = sd(error_score))

(sim_discrim_df_sec_summ)

## # A tibble: 16 x 5
## # Groups:   model_type, df_type [4]
##   model_type df_type n_obs mean_err sd_err
##   <chr>      <chr>   <fct>   <dbl>  <dbl>
## 1 lda       test     100     0.286 0.0806
## 2 lda       test    1000     0.274 0.0253
## 3 lda       test   10000     0.273 0.00860
## 4 lda       test  1e+05     0.273 0.00269
## 5 lda       train    100     0.263 0.0531
## 6 lda       train   1000     0.274 0.0171
## 7 lda       train  10000     0.274 0.00548
## 8 lda       train  1e+05     0.273 0.00176
## 9 qda       test     100     0.273 0.0807
## 10 qda      test    1000     0.261 0.0249
## 11 qda      test   10000     0.261 0.00821
## 12 qda      test  1e+05     0.261 0.00258
## 13 qda      train    100     0.241 0.0475
## 14 qda      train   1000     0.260 0.0160
## 15 qda      train  10000     0.261 0.00525
## 16 qda      train  1e+05     0.261 0.00168

sim_discrim_df_sec_summ %>%
  ggplot(aes(x = n_obs, y = mean_err, group = model_type, colour = model_type)) +
  geom_line() +
  geom_point() +
  geom_errorbar(aes(ymin=mean_err-sd_err, ymax=mean_err+sd_err), width=.2,
    position=position_dodge(0.05)) +
  facet_wrap(vars(df_type)) + theme_minimal() +
  ggtitle("Mean Error Rate (with error bars) of QDA vs LDA over multiple sample sizes")

```



**Conclusion:** We notice that QDA has lower error scores relative to LDA consistently, but the effect of increasing sample size is the lower variance in mean error. We get much more stable estimates, with almost no overlap in the ranges at  $n = 10000$ .

## Q5 - Mental Health Data

```
mh_df <- read_csv("../problem-set-2/mental_health.csv") %>%
  mutate(age = as.integer(age),
         black = as_factor(black),
         educ = as.integer(educ),
         female = as_factor(female),
         married = as_factor(married),
         mhealth_sum = as.integer(mhealth_sum),
         vote96 = as_factor(vote96))
#head(mh_df)
summary(mh_df)
```

```
##   vote96      mhealth_sum      age      educ      black      female
## 0      : 830   Min.      : 0.000   Min.    :18.00   Min.      : 0.00   0:2432   0:1232
## 1     :1783   1st Qu.: 1.000   1st Qu.:32.00   1st Qu.:12.00   1: 400   1:1600
## NA's: 219   Median : 2.000   Median :42.00   Median :13.00
##              Mean    : 2.869   Mean    :45.56   Mean     :13.25
##              3rd Qu.: 4.000   3rd Qu.:57.00   3rd Qu.:16.00
##              Max.     :16.000   Max.     :89.00   Max.      :20.00
##              NA's     :1418    NA's      :4      NA's       :12
## married      inc10
## 0      :1485   Min.      : 0.0535
## 1      :1346   1st Qu.: 2.0062
## NA's:    1   Median : 3.4774
##              Mean     : 4.5761
##              3rd Qu.: 5.8849
##              Max.      :14.8796
##              NA's      :329
```

We have a large number of NAs, and I don't have a good theory for imputation, so I'll go ahead and drop them all for this analysis.

Additionally, the HW prompt says that mhealth\_sum has values from 0 to 9. So I'll drop anything greater than 9.

```
mh_df_clean <- mh_df %>% na.omit() %>% filter(mhealth_sum <= 9)

mh_df_split <- rsample::initial_split(mh_df_clean, prop = 0.7)
mh_df_train <- rsample::training(mh_df_split)
mh_df_test <- rsample::testing(mh_df_split)

logit_mh_model <- glm(vote96 ~ ., family = "binomial", data = mh_df_train)
lda_mh_model <- MASS::lda(vote96 ~ ., data = mh_df_train)
qda_mh_model <- MASS::qda(vote96 ~ ., data = mh_df_train)
nb_mh_model <- train(x = mh_df_train %>% select(-vote96), y = mh_df_train$vote96, method = "nb")

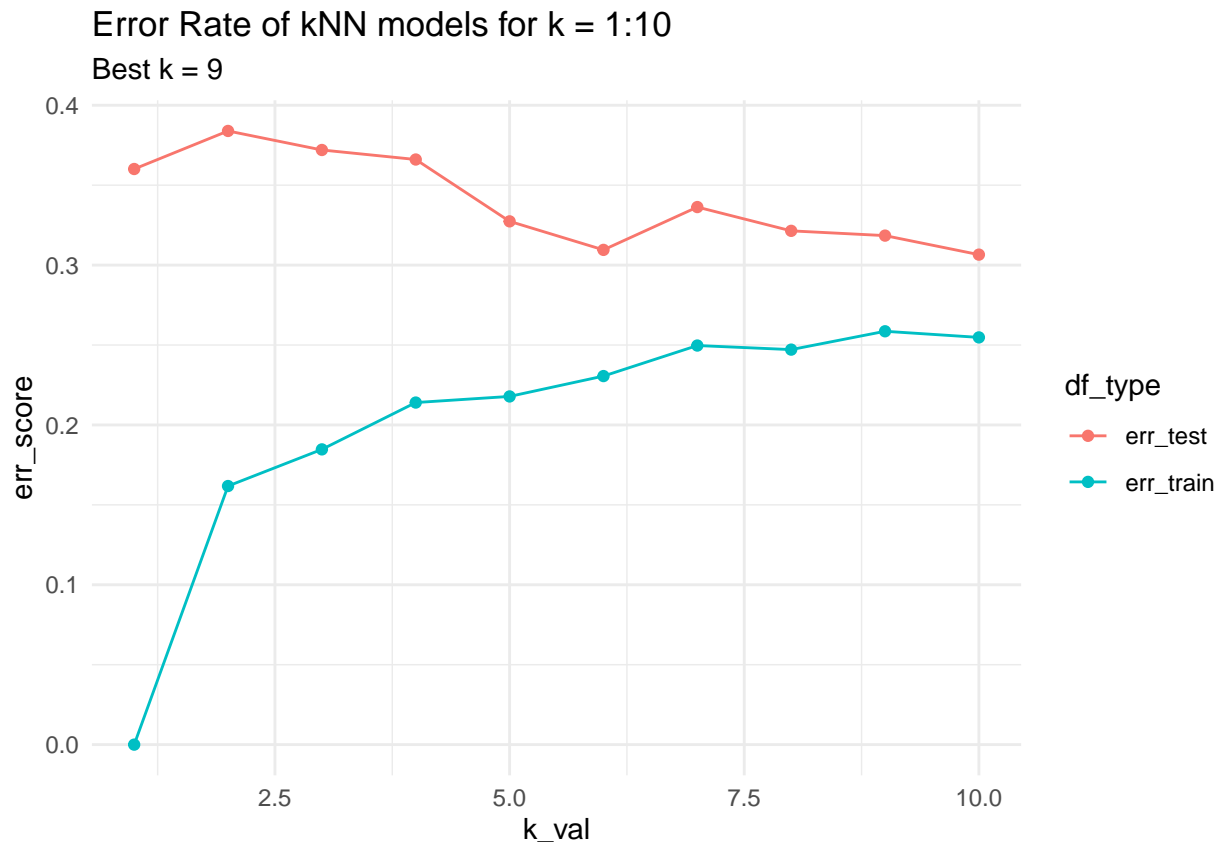
knn_mh_model <- tibble(k_val = 1:10,
                      knn_mod = map(k_val, ~knn3( formula = vote96 ~ .,
                                                    data=mh_df_train, k=.x))) %>%
```

```

mutate(train_classes = map(knn_mod, ~predict(.x, newdata = mh_df_train, type = "class")),
       test_classes = map(knn_mod, ~predict(.x, newdata = mh_df_test, type = "class")),
       err_train = 1 - map_dbl(train_classes, ~postResample(.x, mh_df_train$vote96)[1]),
       err_test = 1 - map_dbl(test_classes, ~postResample(.x, mh_df_test$vote96)[1]))

knn_mh_model %>% pivot_longer(c("err_test", "err_train"), names_to = "df_type", values_to = "err_score")
ggplot(aes(x = k_val, y = err_score, colour = df_type)) +
  geom_line() + geom_point() + theme_minimal() +
  ggtitle("Error Rate of kNN models for k = 1:10", subtitle = "Best k = 9")

```



```

knn_mh_model_best <- knn_mh_model[9, 'knn_mod']$knn_mod[[1]]

logit_res <- predict(logit_mh_model, mh_df_test)
test_prob <- tibble(
  logit = exp(logit_res)/(1+exp(logit_res)),
  lda = predict(lda_mh_model, mh_df_test)$posterior[, '1'],
  qda = predict(qda_mh_model, mh_df_test)$posterior[, '1'],
  nbayes = predict(nbayes_mh_model, mh_df_test, type = 'prob')$'1',
  knn_9 = predict(knn_mh_model_best, newdata = mh_df_test, type = "prob")[, '1'])

mh_df_test_pred <- mh_df_test %>% bind_cols(test_prob %>% mutate_all(~if_else(.x > 0.5, 1, 0)))

class_err_test <- tibble(
  logit = 1 - postResample(mh_df_test_pred$logit, mh_df_test_pred$vote96)[1],
  lda = 1 - postResample(mh_df_test_pred$lda, mh_df_test_pred$vote96)[1],
  qda = 1 - postResample(mh_df_test_pred$qda, mh_df_test_pred$vote96)[1],

```

```

    nbayes = 1 - postResample(mh_df_test_pred$nbayes, mh_df_test_pred$vote96)[1],
    knn_9 = 1 - postResample(mh_df_test_pred$knn_9, mh_df_test_pred$vote96)[1]
  )

class_err_test

## # A tibble: 1 x 5
##   logit   lda   qda nbayes knn_9
##   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.277 0.280 0.310 0.286 0.318

test_prob_roc <- lapply(test_prob, pROC::roc, response = mh_df_test$vote96)
test_auc <- lapply(test_prob_roc, pROC::auc)
test_auc

## $logit
## Area under the curve: 0.7674
##
## $lda
## Area under the curve: 0.7658
##
## $qda
## Area under the curve: 0.7425
##
## $nbayes
## Area under the curve: 0.7604
##
## $knn_9
## Area under the curve: 0.7086

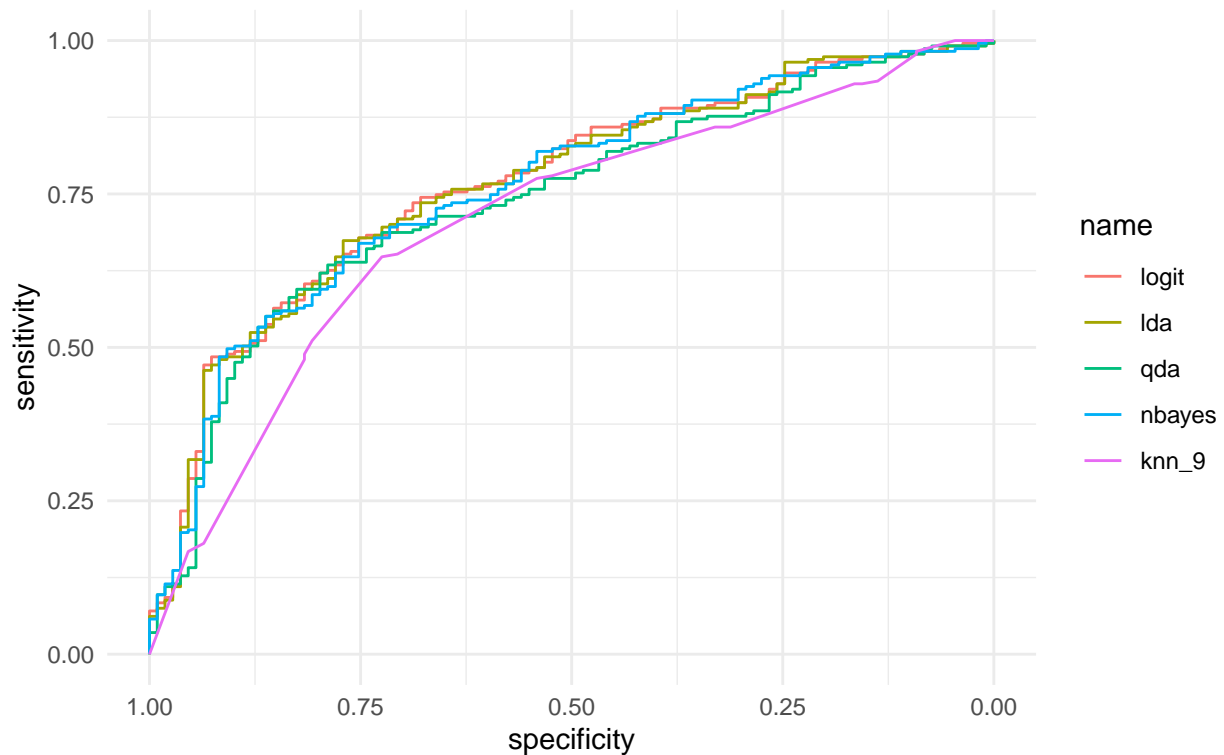
ggroc(test_prob_roc) +
  theme_minimal() +
  ggtitle("ROC Curves", subtitle = "Logit and LDA have the highest AUC")

```



## ROC Curves

Logit and LDA have the highest AUC



**Conclusion:** We have two metrics of identifying the 'best' model here: AUC and `test_error_rate`.

- Naive Bayes has the lowest test error rate (0.2857143), while its AUC is 0.7337.
- Logit has a higher test error rate (0.2767857), but it has the highest AUC of 0.745.

The two models are largely comparable without much relative loss in performance. Since this is a binary classification problem, our choice of threshold (0.5 here) plays a key role in class assignment. Our accuracy measures might waver as we change the threshold. AUC is a more stable metric for binary classification. So, if forced to make a choice, I would choose the Logit model.